

seoul-bike-pranali

May 20, 2024

1 Project Name Machine Learning Project on Bike Rent Demand

Project Type - Regression Pranali Yadav

2 Project Summary -

Bike Seoul is a bike sharing service in the city of Seoul. It is part of the city's efforts to promote sustainable transportation and reduce traffic congestion. The service allows residents and visitors to rent bicycles at various stations across the city and return them to any other station, providing a convenient and eco-friendly mode of transportation. In recent years, the demand for bike rentals in Seoul has increased, leading to the need for a more efficient and effective way to manage the bike sharing operations. Accurately predicting bike demand is crucial for optimizing fleet management, ensuring the availability of bikes at high-demand locations, and reducing waste and costs.

The main objective of this project is to develop a machine learning model that can accurately predict the demand for bike rentals in Seoul, South Korea, based on historical data and various relevant factors such as weather conditions, time of day, and public holidays. In this project we have used regression analysis techniques to model the bike demand data. The model trained on a large dataset of past bike rental information, along with relevant weather and time data. The model then be tested and evaluated using metrics such as mean squared error and r-squared values. The actual data is from the Seoul city government's open data portal, and this dataset is also available on Kaggle.

So, our main goal was to achieve an accuracy of at least 85% in the bike demand predictions, which would help the city's bike sharing service providers plan their fleet operations more effectively and respond to demand changes in real-time. We have performed lots of regression algorithms like linear regression, random forest, decision tree, gradient boosting, Xtreme gradient boosting, also we tried to do hyperparameter tuning and cross validation to improve the accuracy of the model. And finally we have decided to select Xtreme gradient boosting algorithm because it gave us high accuracy around 93% and 90% on train and test data respectively.

This project not only provided valuable insights into bike demand patterns in Seoul but also demonstrated the practical applications of machine learning in addressing real-world problems. The findings could potentially be extended to other cities with similar bike sharing systems, leading to improved services for bike users and more sustainable transportation systems.

3 Problem Statement

Currently Rental bikes are introduced in many urban cities. The business problem is to ensure a stable supply of rental bikes in urban cities by predicting the demand for bikes at each hour. By providing a stable supply of rental bikes, the system can enhance mobility comfort for the public and reduce waiting time, leading to greater customer satisfaction.

To address this problem, i need to develop a predictive model that takes into account various factors that may influence demand, such as time of day, seasonality, weather conditions, and holidays. By accurately predicting demand, the bike sharing system operators can ensure that there is an adequate supply of bikes available at all times, which can improve the user experience and increase usage of the bike sharing system. This can have a positive impact on the sustainability of urban transportation, as it can reduce congestion, air pollution, and greenhouse gas emissions.

```
[5]: # Import Libraries
import numpy as np
import pandas as pd
```

```
[6]: from numpy import math

from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
#from xgboost import XGBRFRegressor

import warnings
warnings.filterwarnings('ignore')
```

```
C:\Users\ADMIN\AppData\Local\Temp\ipykernel_10804\2533403659.py:1:
DeprecationWarning: `np.math` is a deprecated alias for the standard library
`math` module (Deprecated Numpy 1.25). Replace usages of `np.math` with `math`
from numpy import math
```

3.0.1 Dataset Loading

```
[7]: df = pd.read_csv('c:\\Users\\ADMIN\\Downloads\\SeoulBikeData.csv',header=0,
    ↳encoding= 'unicode_escape')
```

3.0.2 Dataset Information

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 8760 entries, 0 to 8759

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	Date	8760 non-null	object
1	Rented Bike Count	8760 non-null	int64
2	Hour	8760 non-null	int64
3	Temperature(°C)	8760 non-null	float64
4	Humidity(%)	8760 non-null	int64
5	Wind speed (m/s)	8760 non-null	float64
6	Visibility (10m)	8760 non-null	int64
7	Dew point temperature(°C)	8760 non-null	float64
8	Solar Radiation (MJ/m2)	8760 non-null	float64
9	Rainfall(mm)	8760 non-null	float64
10	Snowfall (cm)	8760 non-null	float64
11	Seasons	8760 non-null	object
12	Holiday	8760 non-null	object
13	Functioning Day	8760 non-null	object

dtypes: float64(6), int64(4), object(4)

memory usage: 958.3+ KB

```
[9]: # Dataset Describe (all columns included)
df.describe()
```

```
[9]:
```

	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	\
count	8760.000000	8760.000000	8760.000000	8760.000000	
mean	704.602055	11.500000	12.882922	58.226256	
std	644.997468	6.922582	11.944825	20.362413	
min	0.000000	0.000000	-17.800000	0.000000	
25%	191.000000	5.750000	3.500000	42.000000	
50%	504.500000	11.500000	13.700000	57.000000	
75%	1065.250000	17.250000	22.500000	74.000000	
max	3556.000000	23.000000	39.400000	98.000000	

	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	\
count	8760.000000	8760.000000	8760.000000	
mean	1.724909	1436.825799	4.073813	
std	1.036300	608.298712	13.060369	
min	0.000000	27.000000	-30.600000	
25%	0.900000	940.000000	-4.700000	
50%	1.500000	1698.000000	5.100000	
75%	2.300000	2000.000000	14.800000	
max	7.400000	2000.000000	27.200000	

	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)
count	8760.000000	8760.000000	8760.000000
mean	0.569111	0.148687	0.075068

std	0.868746	1.128193	0.436746
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.010000	0.000000	0.000000
75%	0.930000	0.000000	0.000000
max	3.520000	35.000000	8.800000

3.0.3 Dataset First View

```
[10]: df.head()
```

```
[10]:
```

	Date	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	\
0	01/12/2017	254	0	-5.2	37	
1	01/12/2017	204	1	-5.5	38	
2	01/12/2017	173	2	-6.0	39	
3	01/12/2017	107	3	-6.2	40	
4	01/12/2017	78	4	-6.0	36	

	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	\
0	2.2	2000	-17.6	
1	0.8	2000	-17.6	
2	1.0	2000	-17.7	
3	0.9	2000	-17.6	
4	2.3	2000	-18.6	

	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday	\
0	0.0	0.0	0.0	Winter	No Holiday	
1	0.0	0.0	0.0	Winter	No Holiday	
2	0.0	0.0	0.0	Winter	No Holiday	
3	0.0	0.0	0.0	Winter	No Holiday	
4	0.0	0.0	0.0	Winter	No Holiday	

	Functioning Day
0	Yes
1	Yes
2	Yes
3	Yes
4	Yes

3.0.4 Dataset Rows & Columns count

```
[11]: # Dataset Rows & Columns count
# Checking number of rows and columns of the dataset using shape
print("Number of rows are: ",df.shape[0])
print("Number of columns are: ",df.shape[1])
```

Number of rows are: 8760

Number of columns are: 14

Missing Values/Null Values

```
[12]: # Checking Missing Values/Null Values Count
df.isnull().sum()
```

```
[12]: Date                                0
Rented Bike Count                        0
Hour                                    0
Temperature(°C)                         0
Humidity(%)                             0
Wind speed (m/s)                        0
Visibility (10m)                        0
Dew point temperature(°C)               0
Solar Radiation (MJ/m2)                 0
Rainfall(mm)                           0
Snowfall (cm)                           0
Seasons                                 0
Holiday                                 0
Functioning Day                         0
dtype: int64
```

```
[13]: df.value_counts(normalize=True, sort=True, dropna=True, ascending=False)
```

```
[13]: Date      Rented Bike Count  Hour  Temperature(°C)  Humidity(%)  Wind speed
(m/s)  Visibility (10m)  Dew point temperature(°C)  Solar Radiation (MJ/m2)
Rainfall(mm)  Snowfall (cm)  Seasons  Holiday      Functioning Day
01/01/2018    61              5      -4.4              54          0.8
1786          -12.3              0.00              0.0
0.0          Winter  Holiday      Yes              0.000114
21/04/2018    707              0      16.3              59          0.2
756           8.2              0.00              0.0
0.0          Spring  No Holiday  Yes              0.000114
324           6.7              7      12.5              68          1.1
457           6.7              0.22              0.0
0.0          Spring  No Holiday  Yes              0.000114
436           7.7              2      14.7              63          1.8
611           7.7              0.00              0.0
0.0          Spring  No Holiday  Yes              0.000114
600           5.3              8      14.6              54          0.9
431           5.3              0.89              0.0
0.0          Spring  No Holiday  Yes              0.000114
...
11/02/2018    112              0      -6.9              36          2.1
2000          -19.5              0.00              0.0
0.0          Winter  No Holiday  Yes              0.000114
103           22      -5.8              57          3.2
```

```

1980          -12.9          0.00          0.0
0.0          Winter    No Holiday    Yes          0.000114
          95          1          -7.0          38          2.6
2000          -18.9          0.00          0.0
0.0          Winter    No Holiday    Yes          0.000114
          93          2          -7.0          42          1.6
2000          -17.7          0.00          0.0
0.0          Winter    No Holiday    Yes          0.000114
31/12/2017    300          15          3.1          25          3.9
2000          -15.0          0.90          0.0
0.0          Winter    Holiday      Yes          0.000114
Name: proportion, Length: 8760, dtype: float64

```

```

[15]: df["Date"] = pd.to_datetime(df["Date"],dayfirst=True)
df.Date.head()

```

```

[15]: 0    2017-12-01
      1    2017-12-01
      2    2017-12-01
      3    2017-12-01
      4    2017-12-01
Name: Date, dtype: datetime64[ns]

```

#Data Visualization

```

[22]: df.corr(numeric_only=True)

```

```

[22]:          Rented Bike Count      Hour  Temperature(°C) \
Rented Bike Count          1.000000  0.410257          0.538558
Hour                  0.410257  1.000000          0.124114
Temperature(°C)       0.538558  0.124114          1.000000
Humidity(%)           -0.199780 -0.241644          0.159371
Wind speed (m/s)       0.121108  0.285197         -0.036252
Visibility (10m)        0.199280  0.098753          0.034794
Dew point temperature(°C) 0.379788  0.003054          0.912798
Solar Radiation (MJ/m2)  0.261837  0.145131          0.353505
Rainfall(mm)           -0.123074  0.008715          0.050282
Snowfall (cm)          -0.141804 -0.021516         -0.218405

          Humidity(%)  Wind speed (m/s)  Visibility (10m) \
Rented Bike Count    -0.199780          0.121108          0.199280
Hour                 -0.241644          0.285197          0.098753
Temperature(°C)       0.159371         -0.036252          0.034794
Humidity(%)           1.000000         -0.336683         -0.543090
Wind speed (m/s)      -0.336683          1.000000          0.171507
Visibility (10m)      -0.543090          0.171507          1.000000
Dew point temperature(°C) 0.536894         -0.176486         -0.176630

```

Solar Radiation (MJ/m2)	-0.461919	0.332274	0.149738
Rainfall(mm)	0.236397	-0.019674	-0.167629
Snowfall (cm)	0.108183	-0.003554	-0.121695

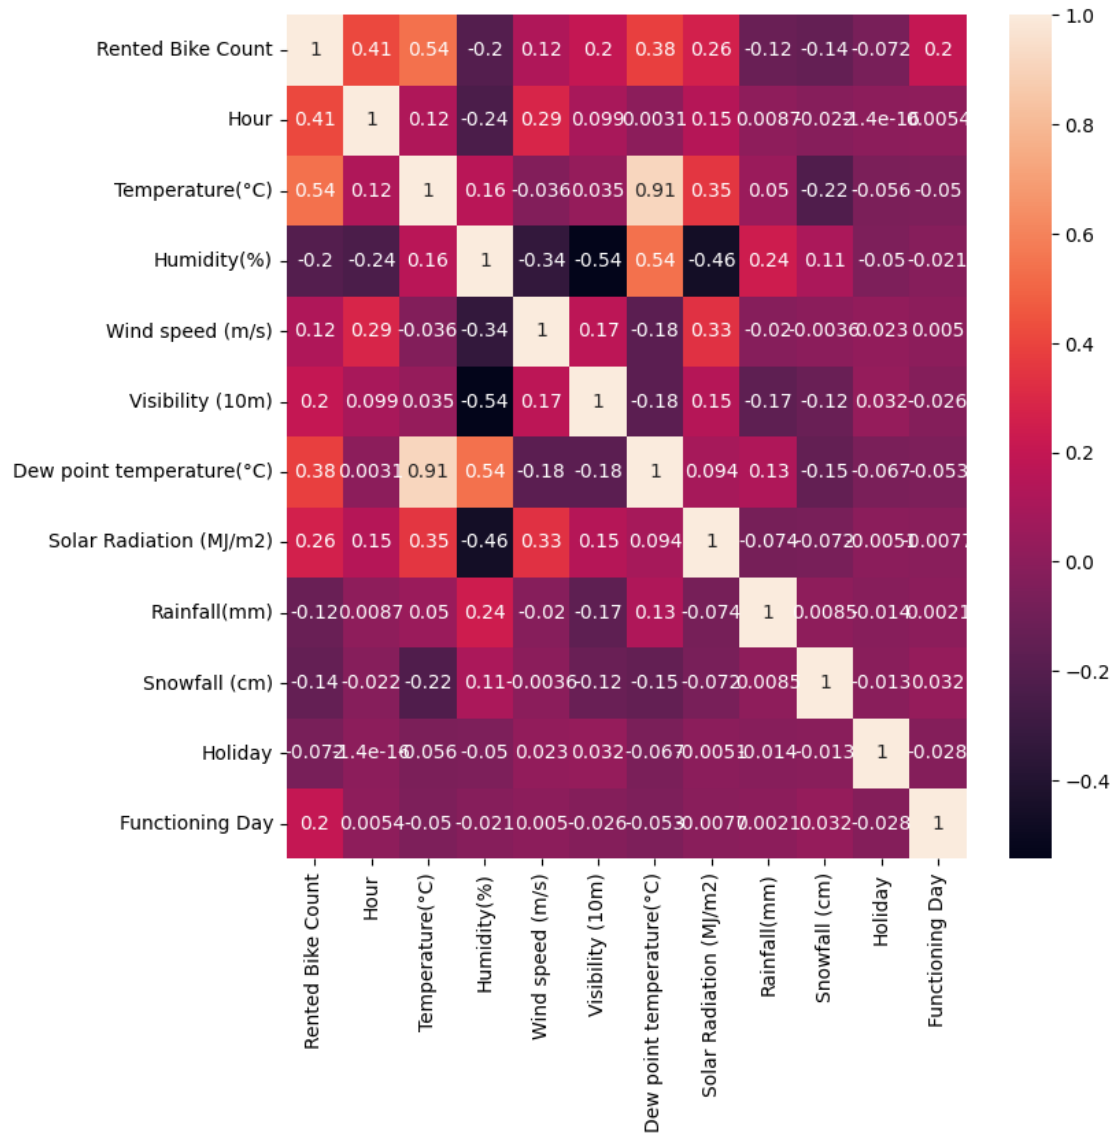
	Dew point temperature(°C)	Solar Radiation (MJ/m2) \
Rented Bike Count	0.379788	0.261837
Hour	0.003054	0.145131
Temperature(°C)	0.912798	0.353505
Humidity(%)	0.536894	-0.461919
Wind speed (m/s)	-0.176486	0.332274
Visibility (10m)	-0.176630	0.149738
Dew point temperature(°C)	1.000000	0.094381
Solar Radiation (MJ/m2)	0.094381	1.000000
Rainfall(mm)	0.125597	-0.074290
Snowfall (cm)	-0.150887	-0.072301

	Rainfall(mm)	Snowfall (cm)
Rented Bike Count	-0.123074	-0.141804
Hour	0.008715	-0.021516
Temperature(°C)	0.050282	-0.218405
Humidity(%)	0.236397	0.108183
Wind speed (m/s)	-0.019674	-0.003554
Visibility (10m)	-0.167629	-0.121695
Dew point temperature(°C)	0.125597	-0.150887
Solar Radiation (MJ/m2)	-0.074290	-0.072301
Rainfall(mm)	1.000000	0.008500
Snowfall (cm)	0.008500	1.000000

```
[78]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8, 8))
sns.heatmap(df.corr(numeric_only=True), annot=True)
```

[78]: <Axes: >




Only Temperature and Hour has considerable correlation with our target variable 'Rented Bike Count'. There is correlation between Temperature and Dew point Temperature(°C).

```
[85]: temperature = df['Temperature(°C)']
rented_bike_count = df['Rented Bike Count']

# Plot the scatter plot
plt.figure(figsize=(4, 4))
plt.scatter(temperature, rented_bike_count, alpha=0.5)
plt.title('Temperature vs. Rented Bike Count')
plt.xlabel('Temperature (°C)')
plt.ylabel('Rented Bike Count')
```

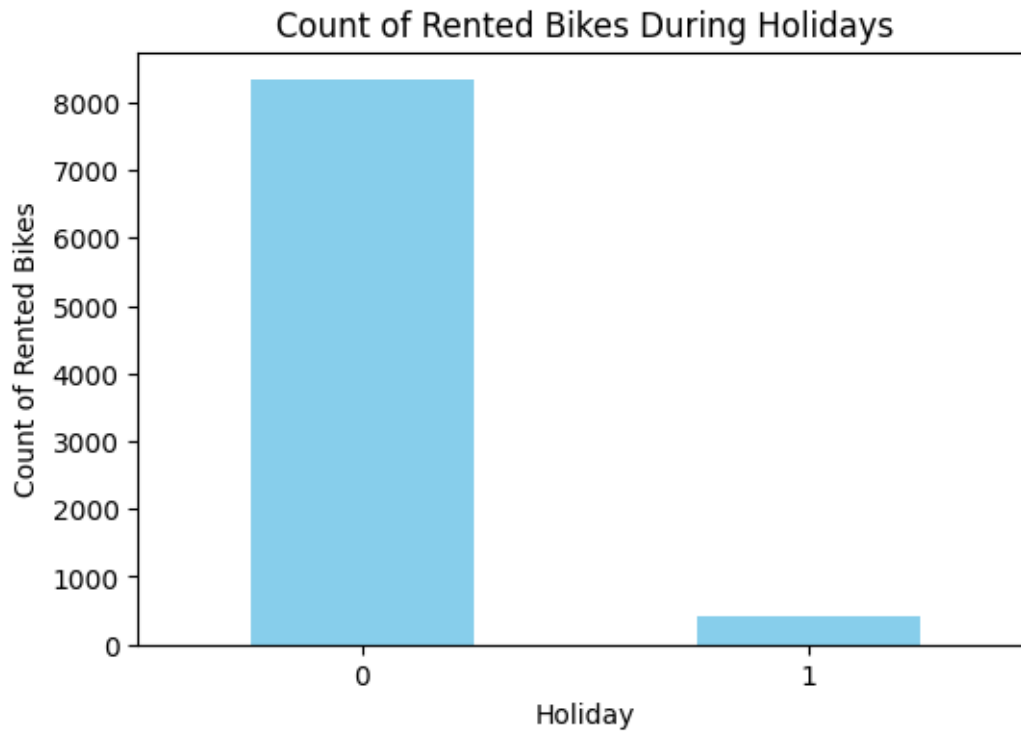


```
plt.show()
```




```
[82]: holiday_counts = df['Holiday'].value_counts()

# Plot the bar graph
plt.figure(figsize=(6, 4))
holiday_counts.plot(kind='bar', color='skyblue')
plt.title('Count of Rented Bikes During Holidays')
plt.xlabel('Holiday')
plt.ylabel('Count of Rented Bikes')
plt.xticks(rotation=0)
plt.show()
```



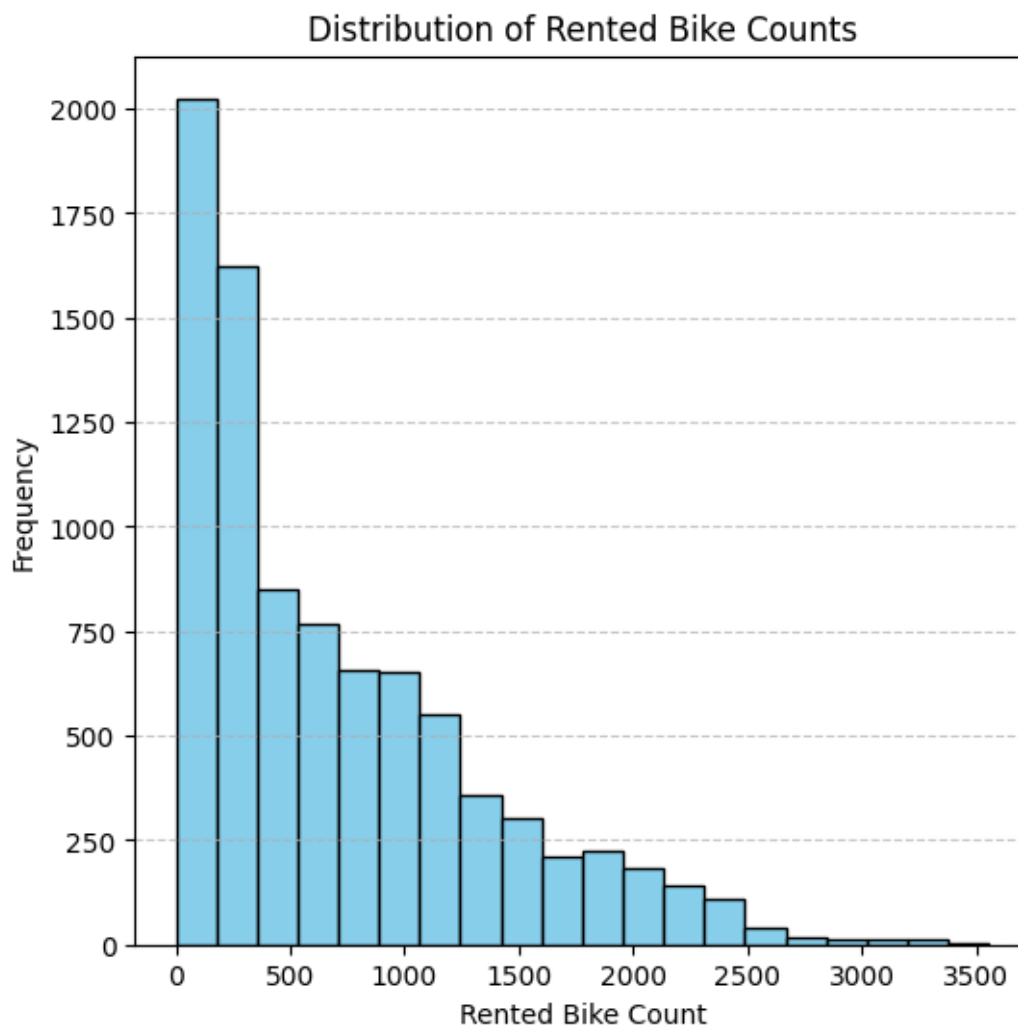
```
[84]: season_counts = df['Seasons'].value_counts()

# Plot the pie chart
plt.figure(figsize=(4, 4))
plt.pie(season_counts, labels=season_counts.index, autopct='%1.1f%%',
        ↪startangle=140)
plt.title('Distribution of Rented Bike Counts by Season')
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle
plt.show()
```



```
[89]: rented_bike_counts = df['Rented Bike Count']

# Plot the histogram
plt.figure(figsize=(6, 6))
plt.hist(rented_bike_counts, bins=20, color='skyblue', edgecolor='black')
plt.title('Distribution of Rented Bike Counts')
plt.xlabel('Rented Bike Count')
plt.ylabel('Frequency')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



Converting categorical data into machine redable data

```
[16]: new_df = df.copy()
```

```
[17]: new_df.head()
```

```
[17]:
```

	Date	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	\
0	2017-12-01	254	0	-5.2	37	
1	2017-12-01	204	1	-5.5	38	
2	2017-12-01	173	2	-6.0	39	
3	2017-12-01	107	3	-6.2	40	
4	2017-12-01	78	4	-6.0	36	

	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	\
0	2.2	2000	-17.6	

1	0.8	2000	-17.6
2	1.0	2000	-17.7
3	0.9	2000	-17.6
4	2.3	2000	-18.6

	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday \
0	0.0	0.0	0.0	Winter	No Holiday
1	0.0	0.0	0.0	Winter	No Holiday
2	0.0	0.0	0.0	Winter	No Holiday
3	0.0	0.0	0.0	Winter	No Holiday
4	0.0	0.0	0.0	Winter	No Holiday

	Functioning Day
0	Yes
1	Yes
2	Yes
3	Yes
4	Yes

We have dichotomous in Holiday and Functioning Day we can convert them into 0 nad 1

```
[18]: df['Holiday'] = df['Holiday'].replace('No Holiday',0)
df['Holiday'] = df['Holiday'].replace('Holiday',1)
```

```
[19]: df['Functioning Day'] = df['Functioning Day'].replace('No',0)
df['Functioning Day'] = df['Functioning Day'].replace('Yes',1)
```

For Seasons we can use one hot key encoding as we have nominal data in this feature

```
[20]: categorical_columns2 = new_df.select_dtypes(['object'])
```

```
[21]: encoded = pd.get_dummies(categorical_columns2, dtype='int64')
```

```
[30]: encoded.head()
```

```
[30]:
```

	Seasons_Autumn	Seasons_Spring	Seasons_Summer	Seasons_Winter \
0	0	0	0	1
1	0	0	0	1
2	0	0	0	1
3	0	0	0	1
4	0	0	0	1

	Holiday_Holiday	Holiday_No Holiday	Functioning Day_No \
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0

	Functioning Day_Yes
0	1
1	1
2	1
3	1
4	1

Dropping unnecessary columns • dropping categorical data which has been converted into numerical data • dropping date • Dew point temperature(°C) because its correlated with Temperature(°C) we can keep only one

```
[22]: new_df.drop(categorical_columns2,axis=1,inplace=True)
```

Let's concatenate our encoded dataframe with dataframe that has only numerical columns

```
[23]: new_df = pd.concat([new_df,encoded],axis = 1)
```

```
[24]: new_df.drop('Date', axis=1, inplace=True)
```

```
[25]: new_df.drop('Dew point temperature(°C)', axis=1, inplace=True)
```

Data Splitting Now our dataset is ready for modelling.

```
[26]: X = new_df.drop(columns='Rented Bike Count')
      y = new_df['Rented Bike Count']
```

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=0)
```

```
[37]: print(f'The shapes of train and test set for independent variables are:
      ↪X_train={X_train.shape}, X_test={X_test.shape}')
      print(f'The shapes of train and test set for dependent variables are:
      ↪y_train={y_train.shape}, y_test={y_test.shape}')
```

The shapes of train and test set for independent variables are: X_train=(7008, 16), X_test=(1752, 16)

The shapes of train and test set for dependent variables are: y_train=(7008,), y_test=(1752,)

```
[28]: X_train.values
```

```
[28]: array([[15. ,  8.2, 62. , ...,  1. ,  0. ,  1. ],
            [18. , 28.4, 57. , ...,  1. ,  0. ,  1. ],
            [11. , 29.9, 57. , ...,  1. ,  0. ,  1. ],
            ...,
            [11. , 25.5, 57. , ...,  1. ,  0. ,  1. ],
            [ 0. ,  8.3, 59. , ...,  1. ,  0. ,  1. ],
```

```
[20. ,  7.1, 83. , ...,  1. ,  0. ,  1. ]])
```

```
[29]: from sklearn.preprocessing import StandardScaler
```

```
[30]: from sklearn.model_selection import train_test_split, GridSearchCV, \
      ↪ RandomizedSearchCV
      from sklearn.metrics import \
      ↪ mean_absolute_error, mean_squared_error, root_mean_squared_error, r2_score, accuracy_score
      from sklearn.ensemble import RandomForestRegressor
      from sklearn.tree import DecisionTreeRegressor
      #from xgboost import XGBRFRegressor

      import warnings
      warnings.filterwarnings('ignore')
```

Data Scaling

```
[31]: scaler=StandardScaler()
      X_train=scaler.fit_transform(X_train)
      X_test=scaler.transform(X_test)
```

```
[32]: #after transformation
      X_train
```

```
[32]: array([[ 0.50774916, -0.38874606,  0.18996015, ...,  0.23202281,
               -0.18254726,  0.18254726],
               [ 0.9408263 ,  1.29850393, -0.05608294, ...,  0.23202281,
               -0.18254726,  0.18254726],
               [-0.06968702,  1.42379477, -0.05608294, ...,  0.23202281,
               -0.18254726,  0.18254726],
               ...,
               [-0.06968702,  1.05627497, -0.05608294, ...,  0.23202281,
               -0.18254726,  0.18254726],
               [-1.65763652, -0.38039334,  0.0423343 , ...,  0.23202281,
               -0.18254726,  0.18254726],
               [ 1.22954439, -0.48062601,  1.22334111, ...,  0.23202281,
               -0.18254726,  0.18254726]])
```

ML Model Implementation

#Decision tree regressor

```
[33]: dt_model = DecisionTreeRegressor()
      dt_model.fit(X_train, y_train)
```

```
[33]: DecisionTreeRegressor()
```

```
[34]: reg = DecisionTreeRegressor(criterion='squared_error', max_leaf_nodes=25,
    ↪ random_state=0)
    reg.fit(X_train, y_train)
```

```
[34]: DecisionTreeRegressor(max_leaf_nodes=25, random_state=0)
```

```
[35]: y_predicted = reg.predict(X_test)
    y_train_prediction = reg.predict(X_train)
```

```
[36]: # test results
    print(mean_squared_error(y_test, y_predicted))
    print(r2_score(y_test, y_predicted))
    print(mean_absolute_error(y_test, y_predicted))
    print(root_mean_squared_error(y_test, y_predicted))
```

```
106123.70920344388
0.7464325412031669
219.78681996507268
325.7663414219521
```

```
[37]: # test results
    print(mean_squared_error(y_train, y_train_prediction))
    print(r2_score(y_train, y_train_prediction))
    print(mean_absolute_error(y_train, y_train_prediction))
    print(root_mean_squared_error(y_train, y_train_prediction))
```

```
96596.14855320382
0.7673744199148671
211.18037322540002
310.7992093831704
```

Let's store metric values of train and test set for later comparisons.

CROSS VALIDATION

```
[38]: param_dict = {"criterion": ['squared_error', 'absolute_error'],
    ↪ "max_depth":
    ↪ [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,21,22,23,24]}
```

```
[39]: grid = GridSearchCV(reg,param_grid=param_dict,cv=10,n_jobs=-1)
```

```
[40]: grid.fit(X_train,y_train)
```

```
[40]: GridSearchCV(cv=10,
    estimator=DecisionTreeRegressor(max_leaf_nodes=25, random_state=0),
    n_jobs=-1,
    param_grid={'criterion': ['squared_error', 'absolute_error'],
    'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
    13, 14, 15, 16, 17, 18, 19, 21, 22, 23,
```


24]})

```
[41]: grid.best_estimator_
```

```
[41]: DecisionTreeRegressor(criterion='absolute_error', max_depth=11,  
                           max_leaf_nodes=25, random_state=0)
```

```
[42]: grid.best_params_
```

```
[42]: {'criterion': 'absolute_error', 'max_depth': 11}
```

```
[43]: grid.best_score_
```

```
[43]: 0.7502315477044867
```

```
[44]: dt_model_tunned = grid.best_estimator_
```

```
[45]: y_train_prediction2 =dt_model_tunned.predict(X_train)  
      y_predicted2 =dt_model_tunned.predict(X_test)
```

```
[46]: # train results  
      print(mean_squared_error(y_train, y_train_prediction2))  
      print(r2_score(y_train, y_train_prediction2))  
      print( mean_absolute_error(y_train, y_train_prediction2))  
      print( root_mean_squared_error(y_train, y_train_prediction2))
```

```
103383.40218321918  
0.7510291635406172  
205.53852739726028  
321.5328944030753
```

```
[47]: # test results  
  
      print(mean_squared_error(y_test, y_predicted2))  
      print(r2_score(y_test, y_predicted2))  
      print(mean_absolute_error(y_test, y_predicted2))  
      print( root_mean_squared_error(y_test, y_predicted2))
```

```
108667.37428652968  
0.7403548164799213  
207.33533105022832  
329.6473483687222
```

#RANDOM REGRESSOR

```
[48]: reg = RandomForestRegressor(criterion='squared_error', max_leaf_nodes=25,  
                                  ↪random_state=0)  
      reg.fit(X_train, y_train)
```

```
[48]: RandomForestRegressor(max_leaf_nodes=25, random_state=0)
```

```
[49]: y_predicted1 = reg.predict(X_test)
      y_train_prediction1 = reg.predict(X_train)
```

```
[50]: #TEST RESULT
      print(mean_squared_error(y_test, y_predicted1))
      print(r2_score(y_test, y_predicted1))
      print(mean_absolute_error(y_test, y_predicted1))
      print(root_mean_squared_error(y_test, y_predicted1))
```

```
85796.80872060338
0.795000769164075
199.7319376231824
292.9109228427704
```

```
[51]: #TRAIN RESULT
      print(mean_squared_error(y_train, y_train_prediction1))
      print(r2_score(y_train, y_train_prediction1))
      print(mean_absolute_error(y_train, y_train_prediction1))
      print(root_mean_squared_error(y_train, y_train_prediction1))
```

```
77600.7997196457
0.8131195568329513
191.77284840191143
278.56920095309476
```

• From evaluation metrics result we can conclude our model is overfitting • We need to do hyperparameter tuning to prevent overfitting Cross-validation using GridSearch

```
[52]: # Setting the parameters to tune

      # Number of trees
      n_estimators = list(np.arange(80,200,20,dtype='int64'))

      # Maximum depth of trees
      max_depth = list(np.arange(12,30,2,dtype='int64'))

      # HYperparameter Grid
      param_dict1 = {'n_estimators' : n_estimators,
                     'max_depth' : max_depth}
```

```
[ ]:
```

```
[53]: rf_random = GridSearchCV(estimator=reg,param_grid= param_dict1,cv = 5,n_jobs=-1)
      rf_random.fit(X_train,y_train)
```

```
[53]: GridSearchCV(cv=5,
                  estimator=RandomForestRegressor(max_leaf_nodes=25, random_state=0),
                  n_jobs=-1,
                  param_grid={'max_depth': [12, 14, 16, 18, 20, 22, 24, 26, 28],
                              'n_estimators': [80, 100, 120, 140, 160, 180]})
```

```
[54]: rf_random.best_score_
```

```
[54]: 0.7991689153496643
```

```
[55]: rf_model_tunned = rf_random.best_estimator_
```

```
[56]: y_train_prediction2 = rf_model_tunned.predict(X_train)
      y_predicted2 = rf_model_tunned.predict(X_test)
```

```
[57]: # train results

print( mean_squared_error(y_train, y_train_prediction2))
print(r2_score(y_train, y_train_prediction2))
print( mean_absolute_error(y_train, y_train_prediction2))
print(root_mean_squared_error(y_train, y_train_prediction2))
```

```
77109.84039903947
0.8143018989704468
191.29197361101134
277.686586638677
```

```
[58]: # test results

print( mean_squared_error(y_test, y_predicted2))
print(r2_score(y_test, y_predicted2))
print( mean_absolute_error(y_test, y_predicted2))
print( root_mean_squared_error(y_test, y_predicted2))
```

```
85144.54841046767
0.7965592520945772
199.32218974995013
291.79538791843106
```

#*XGBOOST REGRESSOR

```
[59]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\admin\appdata\local\packages\pythonsoftwarefoundation.python.3.11_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (2.0.3)
Requirement already satisfied: numpy in c:\users\admin\appdata\local\packages\pythonsoftwarefoundation.python.3.11_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from xgboost) (1.26.4)
```

Requirement already satisfied: scipy in c:\users\admin\appdata\local\packages\pythonsoftwarefoundation.python.3.11_qbz5n2kfra8p0\localcache\local-packages\python311\site-packages (from xgboost) (1.13.0)
Note: you may need to restart the kernel to use updated packages.

```
[60]: from xgboost import XGBRFRegressor
```

```
[61]: xgb_model=XGBRFRegressor(random_state=0,n_jobs=-1)
xgb_model.fit(X_train,y_train)
```

```
[61]: XGBRFRegressor(base_score=None, booster=None, callbacks=None,
                    colsample_bylevel=None, colsample_bytree=None, device=None,
                    early_stopping_rounds=None, enable_categorical=False,
                    eval_metric=None, feature_types=None, gamma=None,
                    grow_policy=None, importance_type=None,
                    interaction_constraints=None, max_bin=None,
                    max_cat_threshold=None, max_cat_to_onehot=None,
                    max_delta_step=None, max_depth=None, max_leaves=None,
                    min_child_weight=None, missing=nan, monotone_constraints=None,
                    multi_strategy=None, n_estimators=None, n_jobs=-1,
                    num_parallel_tree=None, objective='reg:squarederror',
                    random_state=0, reg_alpha=None, ...)
```

```
[62]: y_predicted2 = reg.predict(X_test)
y_train_prediction2 = reg.predict(X_train)
```

```
[63]: print(mean_squared_error(y_test, y_predicted2))
print(r2_score(y_test, y_predicted2))
print(mean_absolute_error(y_test, y_predicted2))
print(root_mean_squared_error(y_test, y_predicted2))
```

```
85796.80872060338
0.795000769164075
199.7319376231824
292.9109228427704
```

```
[64]: print(mean_squared_error(y_train, y_train_prediction2))
print(r2_score(y_train, y_train_prediction2))
print(mean_absolute_error(y_train, y_train_prediction2))
print(root_mean_squared_error(y_train, y_train_prediction2))
```

```
77600.7997196457
0.8131195568329513
191.77284840191143
278.56920095309476
```

From evaluation metrics result we can conclude our model is underfitting • We need to do hyper-parameter tuning to prevent underfitting Cross-validation using GridSearch

```
[65]: # Number of trees
n_estimators = list(np.arange(5,15,2,dtype='int64'))

# Maximum depth of trees
max_depth = list(np.arange(10,20,1,dtype='int64'))

# learning rate
learning_rate=list(np.arange(0.05,0.15,0.01))

# gamma
gamma=list(np.linspace(10,20,num=20))

# subsamples
subsample=[0.3,0.5,0.6]

# Hyperparameter Grid
param_dict1 = {'n_estimators' : n_estimators,
               'max_depth' : max_depth,
               'gamma':gamma,
               'subsample':subsample,
               'learning_rate':learning_rate}

[66]: xgb_random = RandomizedSearchCV(estimator=xgb_model,param_distributions =_
    ↪ param_dict1,cv = 5,n_jobs=-1,random_state=12)
xgb_random.fit(X_train,y_train)

[66]: RandomizedSearchCV(cv=5,
                        estimator=XGBRFRegressor(base_score=None, booster=None,
                                                  callbacks=None,
                                                  colsample_bylevel=None,
                                                  colsample_bytree=None, device=None,
                                                  early_stopping_rounds=None,
                                                  enable_categorical=False,
                                                  eval_metric=None,
                                                  feature_types=None, gamma=None,
                                                  grow_policy=None,
                                                  importance_type=None,
                                                  interaction_constraints=None,
                                                  max_bin=None,
                                                  max_cat_threshold=Non...
                                                  17.36842105263158,
                                                  17.894736842105264,
                                                  18.421052631578945,
                                                  18.94736842105263,
                                                  19.473684210526315, 20.0],
                        'learning_rate': [0.05,
                                           0.060000000000000005,
```

```

0.07,
0.080000000000000002,
0.090000000000000001,
0.1,
0.110000000000000001,
0.120000000000000001,
0.13, 0.14],
'max_depth': [10, 11, 12, 13, 14, 15,
               16, 17, 18, 19],
'n_estimators': [5, 7, 9, 11, 13],
'subsample': [0.3, 0.5, 0.6]},

random_state=12)

```

```
[67]: xgb_random.best_estimator_
```

```
[67]: XGBRFRegressor(base_score=None, booster=None, callbacks=None,
                    colsample_bylevel=None, colsample_bytree=None, device=None,
                    early_stopping_rounds=None, enable_categorical=False,
                    eval_metric=None, feature_types=None, gamma=12.105263157894736,
                    grow_policy=None, importance_type=None,
                    interaction_constraints=None, learning_rate=0.14, max_bin=None,
                    max_cat_threshold=None, max_cat_to_onehot=None,
                    max_delta_step=None, max_depth=13, max_leaves=None,
                    min_child_weight=None, missing=nan, monotone_constraints=None,
                    multi_strategy=None, n_estimators=13, n_jobs=-1,
                    num_parallel_tree=None, objective='reg:squarederror',
                    random_state=0, ...)
```

```
[68]: xgb_random.best_score_
```

```
[68]: 0.22151846133825778
```

```
[69]: xgb_model_tunned = xgb_random.best_estimator_
```

```
[72]: y_train_prediction2 = xgb_model_tunned.predict(X_train)
      y_predicted2 = xgb_model_tunned.predict(X_test)
```

```
[75]: # train results

print( mean_squared_error(y_train, y_train_prediction2))
print(r2_score(y_train, y_train_prediction2))
print( mean_absolute_error(y_train, y_train_prediction2))
print( root_mean_squared_error(y_train, y_train_prediction2))
```

```

315844.0674950709
0.23937537347029325
451.8844554500493

```

562.0000600489923

```
[76]: # test results

print( mean_squared_error(y_test, y_predicted2))
print(r2_score(y_test, y_predicted2))
print( mean_absolute_error(y_test, y_predicted2))
print(root_mean_squared_error(y_test, y_predicted2))
```

324480.7585626313

0.22469953232161322

458.67157449243274

569.6321256413048

Conclusion

After having experimented two machine learning algorithms namely; Decision Tree, Random forest and Xgboost, we can say that the best performance is given by Random Fprest model with r2_score of 0.81 and 0.79 for training and test set respectively using RandomSearch. The best hyperparameter values are:

We have come to an end of our long exercise. Throughout the analysis we went through various steps to determine our predictions for bike rent demand. We started with simple EDA where we analysed our dependent variable as well as other independent variables. We found out the correlation, count, relationships with the dependent variable. We looked for missing values and outliers and did some feature modifications.

Finally we implemented 3 machine learning algorithms namely; DecisionTree, RandomForest and XGBoost. We tried hyperparameter tuning to reduce overfitting and increase model performance. The best performance was given by our RandomForest model.

The r2_score of our best model was 0.81 and 0.79 for training and test set respectively. Performance can be improved even further by applying fine tunings and gathering more amount of observations so that the models can identify more patterns and become less prone to overfitting. With evolution of new technology, these numbers can change in future hence there will always be a need to check on the model from time to time. I hope this exercise will help you to take a step forward!